

Module 5

Enforcing NFRs on the Level of API Invocations Using Anypoint API Manager



Objectives

- Describe how **API Manager** controls API invocations
- Use **API policies** to enforce non-functional constraints on API invocations
- Choose between **enforcement of API policies** in an API implementation and an **API proxy**
- Register an **API client** for access to an API version
- Describe when and how to pass **client ID/secret** to an API
- Establish **guidelines for API policies** suitable for System APIs, Process APIs, and Experience APIs


Addressing the NFRs of the "Aggregator Integration" product



NFRs for "Create quote for aggregators"



- Synchronous creation of up to **5 quotes**:
 - Aggregator-defined **XML**-formatted policy description in HTTP POST request
 - **Up to 5 quotes** in Aggregator-defined XML format in HTTP response
- **Performance**:
 - Throughput: up to **1000 requs/s**
 - Response time: median = **200 ms**, maximum = **500 ms** at 1000 requs/s
- **Security: HTTPS mutual authentication**
- **Reliability**: quotes are legally binding and **must not be lost**



Meeting NFRs for "Create quote for aggregators" using Anypoint Platform




- **Throughput and response time:**
 - Must be broken-down to APIs in **all tiers**
 - Must be **enforced, monitored and analyzed**
 - API Manager, Anypoint Analytics
 - Anticipate **caching**
 - Highly **performant** runtime plane for API implementations: **CloudHub**
 - Need to carefully manage **load on Policy Admin System**: API Manager
- Must not lose quotes:
 - Synchronous invocations incl. ACID operation on **Policy Admin System**
- HTTPS mutual authentication:
 - **CloudHub Dedicated Load Balancer**
- Should add **client authentication** on top of HTTPS mutual auth

All contents © MuleSoft Inc.

5


Addressing the NFRs of the "Customer Self-Service App" product



NFRs for "Retrieve policy holder summary"



- Part of "Customer Self-Service App" product
 - Might be opened-up to **external API consumers**
- **Synchronous** HTTP request-response chain
- **Performance:**
 - Ill-defined, aim for **100 requs/s**
 - Aim for avg response time of **2 s** at 100 requs/s
- **Consistency:**
 - Include claims from "Submit auto claim"
- **OAuth 2.0 (customer)**
- **HTTPS**



Meeting the NFRs for "Retrieve policy holder summary" using Anypoint Platform

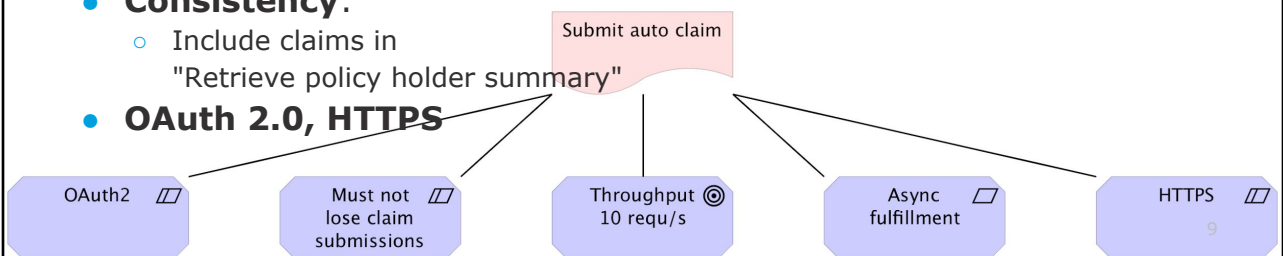


- **Throughput and response time:**
 - **Not challenging**
 - Future use may change that
 - Highly **scalable** runtime plane: **CloudHub**
- **HTTPS:**
 - **Document** in RAML definition
 - Ensure in **API implementation**
- **OAuth 2.0:**
 - Enforce with **API Manager**
 - Requires Identity Provider for **Client Management**
 - PingFederate
- **Consistency:**
 - Through **event notifications**

NFRs for "Submit auto claim"



- Request over HTTP with claim submission and **asynchronous processing** of the submission
 - Processing submission requires lengthy downstream processing steps
- **Performance:**
 - Ill-defined, aim for **10 requs/s**
 - **No response time requirement** because processing is asynchronous
- Reliability: claim submissions **must not be lost**
- **Consistency:**
 - Include claims in "Retrieve policy holder summary"
- **OAuth 2.0, HTTPS**



Meeting the NFRs for "Submit auto claim" using Anypoint Platform



New NFRs for this feature:

- **Async processing** of claim submission and no claim submission loss:
 - **Messaging system**
 - To trigger **async processing without message loss**
 - **Anypoint MQ**
 - Mule runtime **persistent VM queues** as in CloudHub
 - **Persistence mechanism**
 - To store async **correlation** information
 - Mule runtime **Object Store** as in CloudHub
- **Consistency:**
 - Through **event notifications**

Using API Manager and API policies to manage API invocations



Reviewing types of APIs

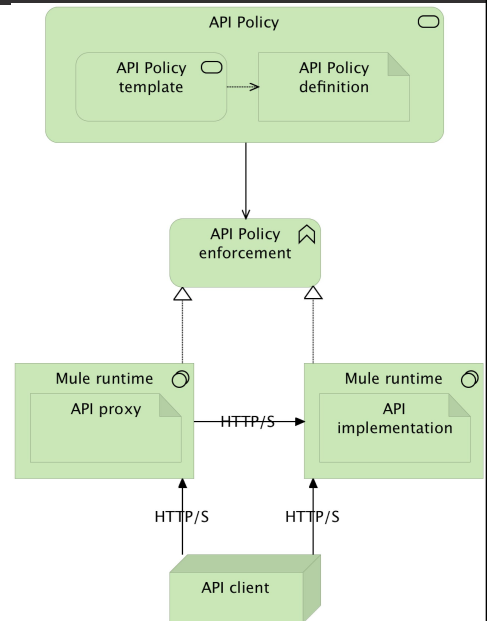


- **REST APIs**
 - With API specification as **RAML** definition or **OAS** definition
 - **Without formal API specification**
 - **Hypermedia**-enabled REST APIs
- **Non-REST APIs**
 - GraphQL APIs
 - SOAP web services (APIs)
 - JSON-RPC, gRPC, ...

- Using **API Manager** and **API policies**
- On the level of **HTTP**
- Applicable to **all types of HTTP/1.x APIs**
 - Therefore not to WebSocket APIs or HTTP/2 APIs
- Special support for **RAML-defined APIs**
 - Allow definition of **resource-level** API policies
 - In addition to the **endpoint-level** API policies available for all APIs

- Defines a typically **non-functional requirement**
- Applied to an **API** (instance)
- Injection into **API invocation** between API client and endpoint
 - Without changing API implementation
- Consists of
 - API policy **template** (code and parameter descriptions)
 - API policy **definition** (parameter values)

- On Anypoint Platform, API policies are always **enforced from within a Mule app**:
 - **API implementation** can **embed** enforcement of API policies
 - **API proxy** deployed in front of the API implementation proper to enforce API policies
- API policies **downloaded at runtime** from API Manager



All contents © MuleSoft Inc.


Exercise: Pros and cons of policy enforcement sites

Compare the characteristics of the two sites of API policies enforcement available in Anypoint Platform:

- List scenarios/requirements that would be best addressed by API policy enforcement **embedded in the API implementation** or in an **API proxy**, respectively

All contents © MuleSoft Inc.

Solution: Pros and cons of policy enforcement sites




- API implementations are **not Mule apps**
- **Resources** must be minized
- **Deployment and CI/CD** must be as simple as possible
- API policies with special **resource requirements** are applied
 - Caching API policy
 - Security API policy requiring HSM
- API policies require **special network configuration**
- **Security sensitive (Experience) APIs**
 - Deployment to **DMZ**
 - **Shield API implementations** from attacks

All contents © MuleSoft Inc.

17

Managing APIs with API Manager



API Administration (Staging)

STAGING

Manage API Search 1 - 6 of 6

API Administration

Client Applications

Custom Policies

Analytics

API Name	Version	Status	Client Applications	Creation Date
Aggregator Quote Creation EAPI	v1	Active	1	01-11-2018 14:21
Home Policy Holder Search SAPI	v1	Active	1	01-11-2018 12:04
Motor Policy Holder Search SAPI	v1	Active	1	01-11-2018 10:28
	v0	Inactive	0	01-11-2018 09:58
Policy Holder Search PAPI				1 version

Aggregator Quote Creation EAPI v1

View API in Exchange

View Analytics Dashboard

Applications Policies SLA tiers

Rate limiting - SLA based

XML threat protection

IP whitelist

Quality of service

Security

Security

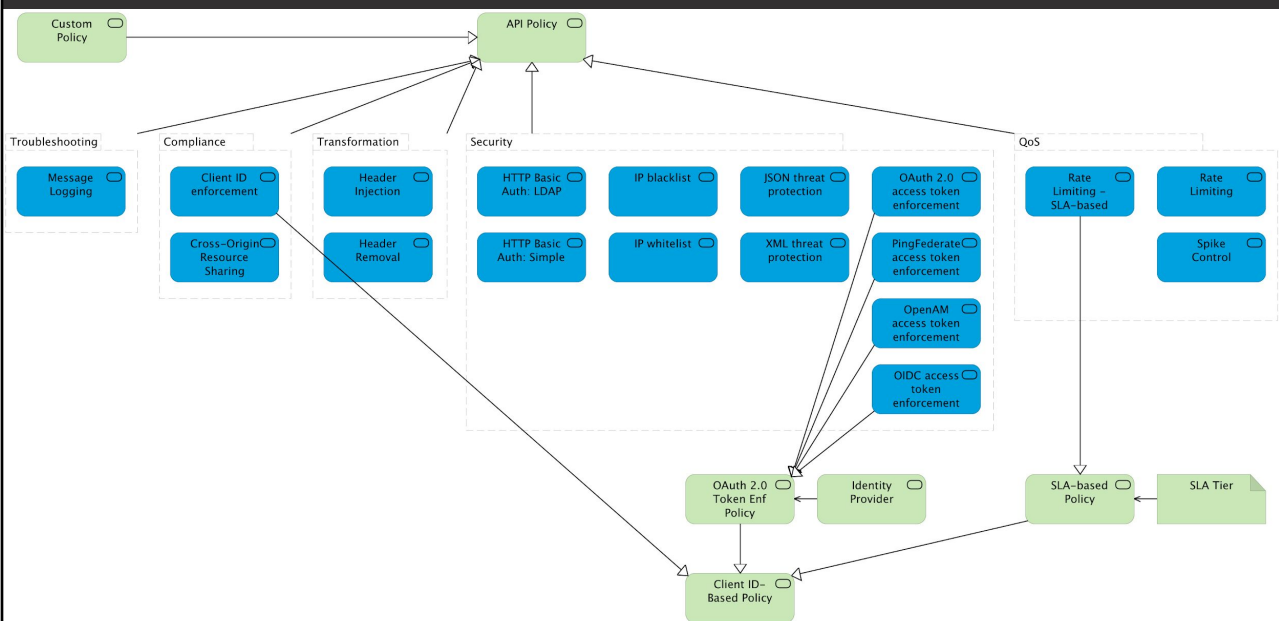
- Management of APIs using **API instances**
 - **API instance** = endpoint for API with major version in environment
- Configuration of **API policies** for a given API instance
 - Select API policy template and parameterize it with API policy definition
 - **OOTB and custom** API policies
- Contacted from site of API policy enforcement to **download all API policies** that must be enforced
- Definition of **alerts** based on API invocations

- Admin of **API clients** ("Client Applications")
 - API consumers use Exchange to request access
- API consumers use **Exchange to request access** to an API
- Access to Anypoint **Analytics**

Selectively applying an API policy to some resources and methods of an API

- By default API policies are applied to **entire API endpoint**
 - Represented as API instance in API Manager
- APIs defined with a RAML definition can apply API policies also to selected combinations of **API resources and HTTP methods**
- **OpenAPI** documents can be converted to **RAML definitions**

API policies



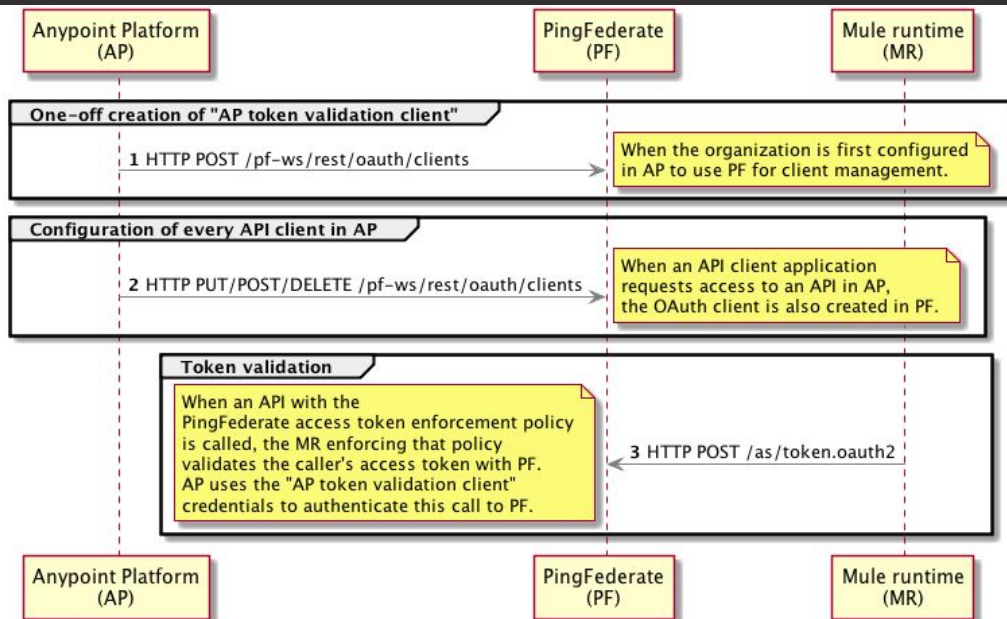
- API policies are **AOP** applied to API invocations:
 - **Ordered**, API implementation/proxy as last element
 - **Incoming HTTP request** passed down this chain, returning **HTTP response** passed up
 - API policies implement "**around advice**":
 - Execute code **before/after** handing control to the **next element** in the chain
 - **Change HTTP request/response** if desired
 - In Mule 4: also applied to **outgoing HTTP requests**

- **Implementing and applying** custom API policies:
 - Very similar to **Mule apps**
 - Packaged and deployed to **Exchange**
 - Contains both **policy template** (code and parameter descriptions)
 - **API Manager** retrieves policy from Exchange and shows **configuration UI** to enter the definition (parameter values)
 - Policy template and definition **downloaded to any Mule runtime** that registers as that API instance

- **Client ID enforcement**
- **CORS control**
 - Interacts with API clients for **Cross-Origin Resource Sharing**:
 - Rejects HTTP requests whose **Origin** request header does not match configured origin domains
 - Sets **Access-Control-*** HTTP response headers to match configured cross-origins, usage of credentials, etc.
 - Responds to CORS pre-flight **HTTP OPTIONS requests**
 - Can be important for Experience APIs invoked from a **browser**

- **Authentication/Authorization**
 - **OAuth 2.0 token enforcement** API policies
 - Require matching Identity Provider configured for **Client Management**
 - OpenAM, PingFederate or OIDC DCR compatible (Okta)
 - Discouraged "OAuth 2.0 access token enforcement using external provider" requires access to Mule OAuth 2.0 provider or other configured in the policy
 - **Basic Authentication: LDAP/Simple**
 - Incorporate access to Identity Provider
- **IP-based** access control
 - **blacklisting, whitelisting**
- **Payload threat protection**
 - Guard against attacks sending over-sized HTTP request bodies
 - **Limit size of XML or JSON bodies**

Interactions with OAuth 2.0 Client Management



All conter

27

QoS-related API policies




- Quality of Service (QoS) related API policies on Anypoint Platform enforce **throughput limit** in # of API invocations per unit of time:
 - **Rate Limiting**: rejects requests above limit
 - **Spike Control**: queues requests above limit
- Two different ways to define the throughput limit:
 - **Non-SLA-based** (Rate Limiting and Spike Control)
 - Limit defined on API policy definition
 - Enforced for that API instance **across all API clients**
 - **SLA-based** (Rate Limiting)
 - Limit defined in an **SLA tier**
 - **API clients must register** with the API instance at a particular SLA tier
 - **Enforced separately** for each registered API client
 - API client must identify itself with **client ID**
- **X-RateLimit-*** HTTP response headers optionally inform API client of remaining capacity

- SLA tiers
 - Enable **different API clients** to receive **different QoS**
 - Define one or more **throughput limits**
 - Per API client and API instance
- API instance with SLA tiers requires every **API client to register** for access with exactly one SLA tier
 - Manual or automatic **approval**
 - API clients must send **client ID/client secret** in API invocations
 - API client is promised the QoS offered by that SLA tier
- Enforcement by **SLA-based Rate Limiting** API policy
- Violation of SLA **monitored, reported and alerted-on**

- API clients must **register to invoke** API instance with Client ID-based API Policies
 - Called "application" or "client application"
- Request access **through Exchange** entry for that API
 - Directly from Exchange or via Public (Developer) Portal
- Access **approval** is automatic or manual
- API consumer receives **client ID and client secret**
 - Must be supplied by that API client in all API invocations to that API version in that environment

Registering API clients with an Anypoint Platform-managed API



Aggregator Quote Creation EAPI v1

Share Download Edit Request access

Request API access

[Create a new application](#)

Application	Aggregator	▼
API Instance	Staging - v1:7484080	▼
SLA tier	Standard	▼
# of Reqs	Time period	Time Unit
1000	1	Second

Cancel

Request API access

Overview


Type REST API
Created By AnySurance Owner
Published On Jan 11, 2018
Visibility Private

Asset versions for v1

Version Instances

1.0.1 Mocking Service
Staging - v1:7484080

Registering API clients with an Anypoint Platform-managed API



API Manager

Acme Insurance ? AO

API Administration (Staging) Aggregator Quote Creation EAPI (v1) - Client Applications

STAGING

Aggregator Quote Creation EAPI v1

Actions

API Status: Active Asset Version: 1.0.1 Type: RAML/OAS
Implementation URL: <http://ans-aggregatorquotecreation-eapi.cloudhub.io/v1>
Consumer endpoint: <http://ans-aggregatorquotecreation-eapi.cloudhub.io/v1>

View API in Exchange
View configuration details
View Analytics Dashboard

Search

1 - 1 of 1

Application	Current SLA tier	Requested SLA tier	Status	
Aggregator	Standard	N/A	Approved	Revoke
Owners	AnySurance Owner anysurance+owner@googlegroups.com		Submitted	5 days ago
Client ID	552f92bfd0a94500b007c165fde8dbd2		Approved	5 days ago
URL	None		Rejected	-
Redirect URIs	None		Revoked	-

- API policies that require **API clients to identify** themselves:
 - **Client ID enforcement**
 - Rate Limiting - **SLA-based**
 - Retrieve SLA tier by client ID
 - Also enforce presence and validity of **client ID** and secret (optional)
 - **OAuth 2.0** access token enforcement
 - Token implicitly carries client ID
 - Policy **exchanges token for client ID** and passes it to SLA-based API policy
- **Client ID and client secret** passed in API invocations as defined by the API policy
 - **Query parameters**
 - Custom request **headers**
 - Standard **Authorization header** as in HTTP Basic Authentication

- To manipulate **HTTP headers** in requests and responses:
 - **Header Injection**
 - Values are **expressions** and hence dynamically evaluated
 - **Header Removal**
- For instance, to propagate transaction IDs as HTTP headers along chains of API invocations

Exercise: Select API policies for all tiers in Acme Insurance's application network



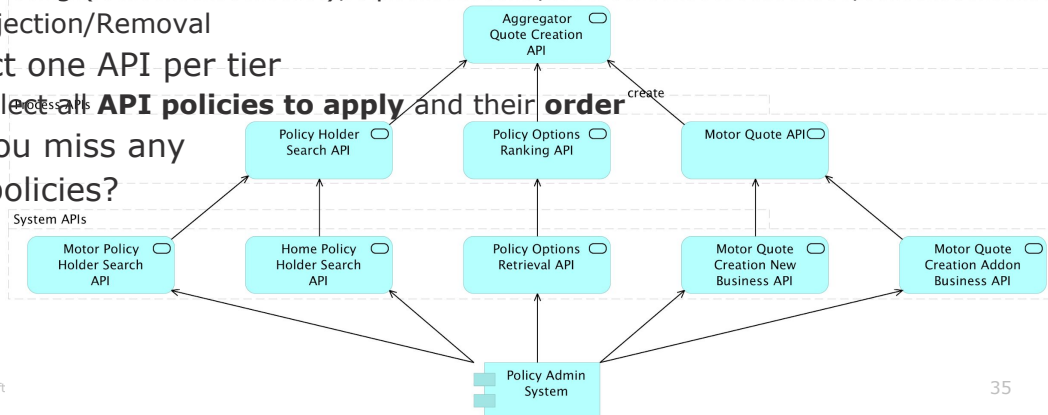
1. Using OOTB API policies

- CORS, HTTP Basic Auth Simple/LDAP, IP black/whitelist, JSON/XML threat protection, PingFederate/OpenAM/OIDC access token enforcement, Rate Limiting (SLA-based or not), Spike Control, Client ID enforcement, Header Injection/Removal

2. Select one API per tier

- Select **all API policies to apply** and their **order**

3. Do you miss any API policies?



All contents © MuleSoft

35

Choosing appropriate API policies for System APIs



Policy Options Retrieval SAPI v1

Actions ▾

API Status: ● Active Asset Version: 1.0.0 Type: RAML/OAS

Implementation URL: <http://ans-policyoptionsretrieval-sapi.cloudhub.io/v1>

Consumer endpoint: <http://ans-policyoptionsretrieval-sapi.cloudhub.io/v1>

[View API in Exchange](#) >

[View configuration details](#) >

[View Analytics Dashboard](#) >

Apply New Policy


Edit policy order

Name	Category	Fulfills	Requires
> IP whitelist ⓘ	Security	IP filtered	
> Rate limiting - SLA based ⓘ	Quality of service	SLA Rate Limiting, Client ID required	RAML snippet
> Spike Control ⓘ	Quality of service	Baseline Rate Limiting	

All contents © MuleSoft Inc.

36

Choosing appropriate API policies for Process APIs



Policy Holder Search PAPI v1

Actions ▾

API Status: ● Active Asset Version: 1.0.3 Type: RAML/OAS

Implementation URL: <http://ans-policyholdersearch-papi.cloudhub.io/v1>

Consumer endpoint: <http://ans-policyholdersearch-papi.cloudhub.io/v1>

View API in Exchange >

View configuration details >

View Analytics Dashboard >

Apply New Policy

Edit policy order

Name	Category	Fulfills	Requires
> IP whitelist ⓘ	Security	IP filtered	
> Client ID enforcement ⓘ	Compliance	Client ID required	RAML snippet
> Spike Control ⓘ	Quality of service	Baseline Rate Limiting	

All contents © MuleSoft Inc.

37

Choosing appropriate API policies for Experience APIs

Aggregator Quote Creation EAPI v1

Actions ▾

API Status: ● Active Asset Version: 1.0.1 Type: RAML/OAS

Implementation URL: <http://ans-aggregatorquotecreation-eapi.cloudhub.io/v1>

Consumer endpoint: <http://ans-aggregatorquotecreation-eapi.cloudhub.io/v1>

View API in Exchange >

View configuration details >

View Analytics Dashboard >

Apply New Policy

Edit policy order

Name	Category	Fulfills	Requires
> IP whitelist ⓘ	Security	IP filtered	
> XML threat protection ⓘ	Security	XML threat protected	
> Rate limiting - SLA based ⓘ	Quality of service	SLA Rate Limiting, Client ID required	RAML snippet

All contents © MuleSoft Inc.

38

Choosing appropriate API policies for Experience APIs

Mobile Policy Holder Summary ... v1

Actions ▾

API Status:  Unregistered Asset Version: 1.0.0 Type: RAML/OAS

Implementation URL: <http://acmeins-mobilepolicyholderssummary-eapi.cloudhub.io/v1>





Consumer endpoint: <http://acmeins-mobilepolicyholderssummary-eapi.cloudhub.io/v1> 

View API in Exchange >

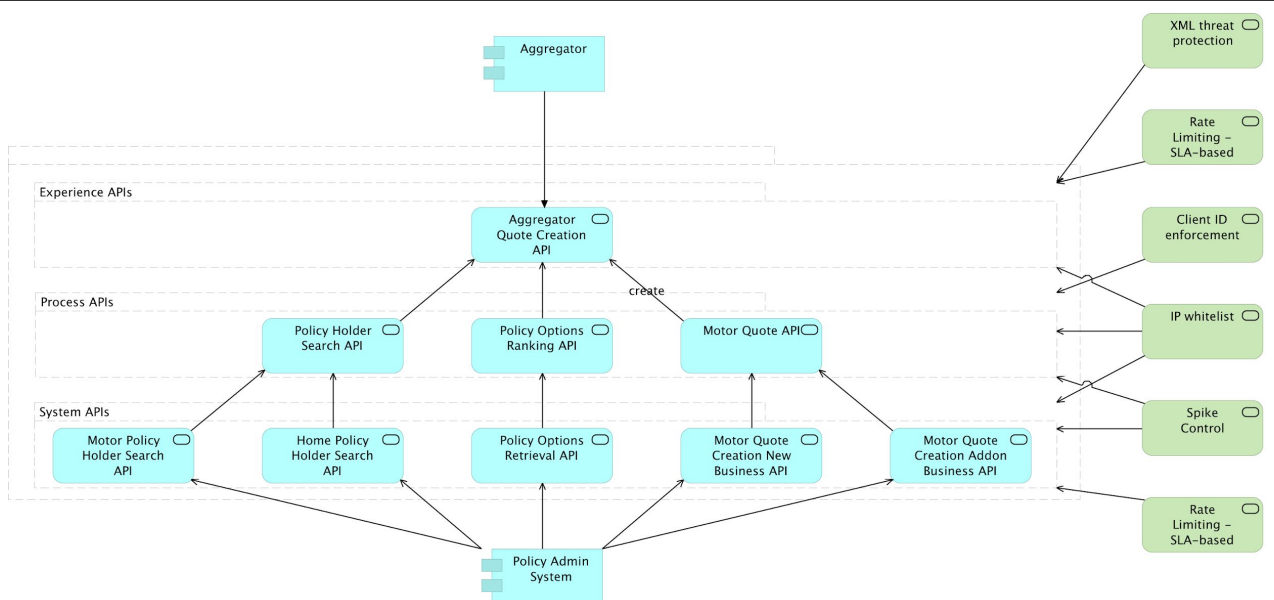
View configuration details >

Apply New Policy


Edit policy order

Name	Category	Fulfills	Requires
> JSON threat protection 	Security	JSON threat protected	
> OAuth 2.0 access token enforcement using external provider 	Security	OAuth 2.0 protected	RAML snippet
> Client ID enforcement 	Compliance	Client ID required	RAML snippet
> Rate limiting 	Quality of service	Baseline Rate Limiting	

API policies for "Create quote for aggregators"



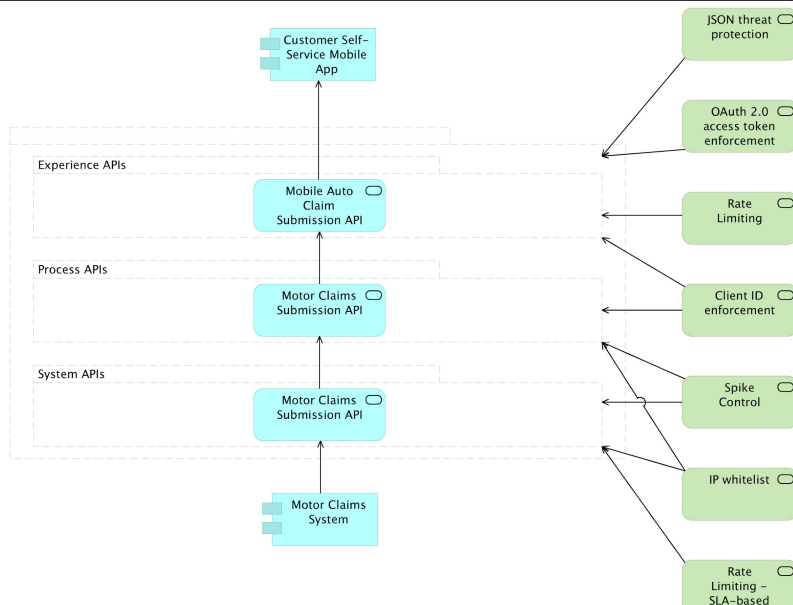
API policies for "Retrieve policy holder summary"



All contents © MuleSoft Inc.

41

API policies for "Submit auto claim"




All contents © MuleSoft Inc.

42

Reflecting the application of API policies in the RAML definition of an API

- Many API policies **change HTTP request/response**:
 - Require certain HTTP request headers: **Authorization**
 - Require certain query parameters: **client_id**
 - Add HTTP response headers: **X-RateLimit-Limit**
- **Change contract** between API client and API implementation
- Must be reflected in **RAML definition** of API
 - RAML has specific support for **securitySchemes** such as OAuth 2.0
 - In other cases define **RAML traits**
- **C4E** owns definition of reusable **RAML fragments**
 - Publish to **Exchange** to encourage consumption and reuse.

Latency overhead of applying API policies



- Increase in HTTP request-response latency
- through API policies
- enforced embedded in API implementation

Summary



Summary



- **NFRs** for products are constraints on throughput, response time, security and reliability
- **API Manager and API policies** control invocations of APIs and impose non-functional constraints
- Compliance, Security, QoS, Transformation
- API policies **enforced**
 - Directly in an **API implementation** that is a Mule app
 - In an **API proxy**
- **Client ID**-based API policies require registered API clients
 - Must pass client ID/secret with every API invocation
- **C4E** defines guidelines for API policies and publishes matching reusable RAML fragments to Exchange